

# Date - 17/10/2023

## Team ID - 721

## Project Title - Customer Churn Prediction

### 1.Import Libraries required to create the Customer Churn Model

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

### 2. Load Churn Prediction Dataset

```
In [2]: data = pd.read_csv('WA_Fn-UseC_-Telco-Customer-Churn.csv')
```

### 3.Exploring Dataset

#### 1. Displaying the top 5 rows

```
In [3]: data.head()
```

```
Out[3]:
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service
1	5575-GNVDE	Male	0	No	No	34	Yes	No
2	3668-QPYBK	Male	0	No	No	2	Yes	No
3	7795-CFOCW	Male	0	No	No	45	No	No phone service
4	9237-HQITU	Female	0	No	No	2	Yes	No

5 rows × 21 columns

#### 2. Displaying the bottom 5 rows

In [4]: `data.tail()`

Out[4]:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines
<b>7038</b>	6840-RESVB	Male	0	Yes	Yes	24	Yes	Y
<b>7039</b>	2234-XADUH	Female	0	Yes	Yes	72	Yes	Y
<b>7040</b>	4801-JAZL	Female	0	Yes	Yes	11	No	No phone service
<b>7041</b>	8361-LTMKD	Male	1	Yes	No	4	Yes	Y
<b>7042</b>	3186-AJIEK	Male	0	No	No	66	Yes	N

5 rows × 21 columns

### 3. Displaying the columns

In [5]: `data.columns`

Out[5]:

```
Index(['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents',
      'tenure', 'PhoneService', 'MultipleLines', 'InternetService',
      'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport',
      'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBilling',
      'PaymentMethod', 'MonthlyCharges', 'TotalCharges', 'Churn'],
      dtype='object')
```

### 4. Finding the shape


In [6]: `data.shape`

Out[6]: (7043, 21)


## 4.Data Visualization

In [7]:

```
plt.figure(figsize=(6, 4))
sns.countplot(data=data, x='gender')
plt.title('Gender Distribution')
plt.xlabel('Gender')
plt.ylabel('Count')
plt.show()
```



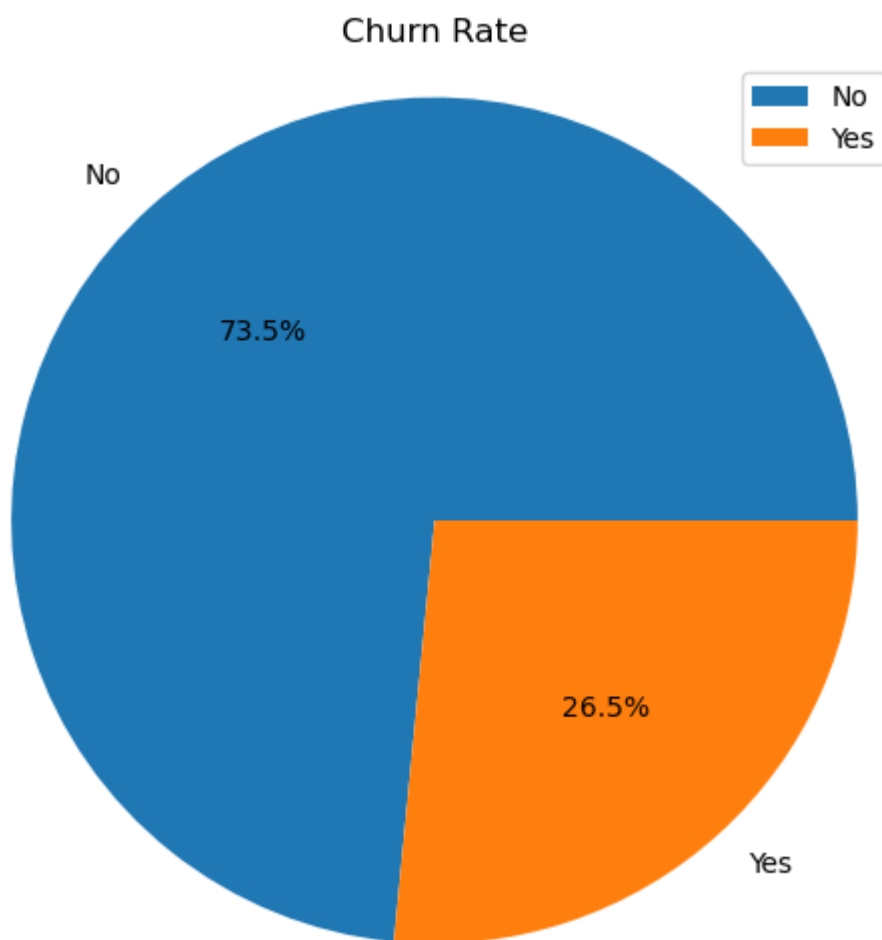
```
In [8]: plt.figure(figsize=(6, 4))
sns.countplot(data=data, x='Contract')
plt.title('Contract Distribution')
plt.xlabel('Contract')
plt.ylabel('Contract(Count)')
plt.show()
```



```
In [9]: # Count the number of customers for each churn category
churn_counts = data['Churn'].value_counts()
```

```
In [10]: # Create a pie chart
plt.figure(figsize=(6, 6))
plt.pie(churn_counts, labels=churn_counts.index, autopct='%1.1f%%')
```

```
plt.title('Churn Rate')  
plt.axis('equal') # Equal aspect ratio ensures that the pie chart is circular.  
plt.legend()  
plt.show()
```



## 5.Preprocess Dataset

```
In [11]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   customerID            7043 non-null   object
1   gender                7043 non-null   object
2   SeniorCitizen         7043 non-null   int64
3   Partner               7043 non-null   object
4   Dependents            7043 non-null   object
5   tenure                7043 non-null   int64
6   PhoneService          7043 non-null   object
7   MultipleLines         7043 non-null   object
8   InternetService       7043 non-null   object
9   OnlineSecurity        7043 non-null   object
10  OnlineBackup          7043 non-null   object
11  DeviceProtection      7043 non-null   object
12  TechSupport           7043 non-null   object
13  StreamingTV           7043 non-null   object
14  StreamingMovies       7043 non-null   object
15  Contract              7043 non-null   object
16  PaperlessBilling      7043 non-null   object
17  PaymentMethod         7043 non-null   object
18  MonthlyCharges        7043 non-null   float64
19  TotalCharges          7043 non-null   object
20  Churn                 7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

```
In [12]: missing_values = data.isnull().sum()
```

```
In [13]: print("Missing Values:\n", missing_values)
```

```
Missing Values:
 customerID            0
gender                0
SeniorCitizen         0
Partner               0
Dependents            0
tenure                0
PhoneService          0
MultipleLines         0
InternetService       0
OnlineSecurity        0
OnlineBackup          0
DeviceProtection      0
TechSupport           0
StreamingTV           0
StreamingMovies       0
Contract              0
PaperlessBilling      0
PaymentMethod         0
MonthlyCharges        0
TotalCharges          0
Churn                 0
dtype: int64
```

```
In [14]: data = data.dropna()
```

```
In [15]: data.describe()
```

Out[15]:

	SeniorCitizen	tenure	MonthlyCharges
count	7043.000000	7043.000000	7043.000000
mean	0.162147	32.371149	64.761692
std	0.368612	24.559481	30.090047
min	0.000000	0.000000	18.250000
25%	0.000000	9.000000	35.500000
50%	0.000000	29.000000	70.350000
75%	0.000000	55.000000	89.850000
max	1.000000	72.000000	118.750000

In [16]:

```
df = data.drop('customerID',axis=1)
```

In [17]:

```
df
```

Out[17]:

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService
0	Female	0	Yes	No	1	No	No phone service	
1	Male	0	No	No	34	Yes	No	
2	Male	0	No	No	2	Yes	No	
3	Male	0	No	No	45	No	No phone service	
4	Female	0	No	No	2	Yes	No	Fiber
...	...	...	...	...	...	...	...	
7038	Male	0	Yes	Yes	24	Yes	Yes	
7039	Female	0	Yes	Yes	72	Yes	Yes	Fiber
7040	Female	0	Yes	Yes	11	No	No phone service	
7041	Male	1	Yes	No	4	Yes	Yes	Fiber
7042	Male	0	No	No	66	Yes	No	Fiber

7043 rows × 20 columns

In [18]:

```
#count of string value into the column.
count=0
for i in df.TotalCharges:
    if i==' ':

```

```

        count+=1
    print('count of empty string:- ',count)
    #we will replace this empty string to nan values
    df['TotalCharges'] = df['TotalCharges'].replace(" ",np.nan)
    # typecasting of the TotalCharges column
    df['TotalCharges'] = df['TotalCharges'].astype(float)

```

count of empty string:- 11

## 6. Checking Null Values in Customer Churn Data

In [19]: `df.isnull().sum()`

```

Out[19]:
gender                0
SeniorCitizen         0
Partner               0
Dependents            0
tenure                0
PhoneService          0
MultipleLines         0
InternetService       0
OnlineSecurity        0
OnlineBackup          0
DeviceProtection      0
TechSupport           0
StreamingTV           0
StreamingMovies       0
Contract              0
PaperlessBilling      0
PaymentMethod         0
MonthlyCharges        0
TotalCharges         11
Churn                 0
dtype: int64

```

In [20]: `# fill null values with mean`  
`df['TotalCharges'] = df['TotalCharges'].fillna(df['TotalCharges'].mean())`

In [21]: `#numerical variables`

```

num = list(df.select_dtypes(include=['int64','float64']).keys())

#categorical variables

cat = list(df.select_dtypes(include='O').keys())

print(cat)

print(num)

```

```

['gender', 'Partner', 'Dependents', 'PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBilling', 'PaymentMethod', 'Churn']
['SeniorCitizen', 'tenure', 'MonthlyCharges', 'TotalCharges']

```

In [22]: `# value_counts of the categorical columns`

```

for i in cat:
    print(df[i].value_counts())
# as we see that there is extra categories which we have to convert it into No.
df.MultipleLines = df.MultipleLines.replace('No phone service','No')
df.OnlineSecurity = df.OnlineSecurity.replace('No internet service','No')
df.OnlineBackup = df.OnlineBackup.replace('No internet service','No')

```

```
df.DeviceProtection = df.DeviceProtection.replace('No internet service','No')  
df.TechSupport = df.TechSupport.replace('No internet service','No')  
df.StreamingTV = df.StreamingTV.replace('No internet service','No')  
df.StreamingMovies = df.StreamingMovies.replace('No internet service','No')
```



```
gender
Male      3555
Female    3488
Name: count, dtype: int64
Partner
No        3641
Yes       3402
Name: count, dtype: int64
Dependents
No        4933
Yes       2110
Name: count, dtype: int64
PhoneService
Yes       6361
No        682
Name: count, dtype: int64
MultipleLines
No        3390
Yes       2971
No phone service    682
Name: count, dtype: int64
InternetService
Fiber optic    3096
DSL            2421
No            1526
Name: count, dtype: int64
OnlineSecurity
No            3498
Yes          2019
No internet service    1526
Name: count, dtype: int64
OnlineBackup
No            3088
Yes          2429
No internet service    1526
Name: count, dtype: int64
DeviceProtection
No            3095
Yes          2422
No internet service    1526
Name: count, dtype: int64
TechSupport
No            3473
Yes          2044
No internet service    1526
Name: count, dtype: int64
StreamingTV
No            2810
Yes          2707
No internet service    1526
Name: count, dtype: int64
StreamingMovies
No            2785
Yes          2732
No internet service    1526
Name: count, dtype: int64
Contract
Month-to-month    3875
Two year          1695
One year          1473
Name: count, dtype: int64
PaperlessBilling
Yes            4171
No            2872
```

```
Name: count, dtype: int64
PaymentMethod
Electronic check      2365
Mailed check          1612
Bank transfer (automatic) 1544
Credit card (automatic) 1522
Name: count, dtype: int64
Churn
No      5174
Yes     1869
Name: count, dtype: int64
```

## 7.Handling categorical Variables in Customer Churn Data

```
In [23]: # we have to handel this all categorical variables
# there are mainly Yes/No features in most of the columns
# we will convert Yes = 1 and No = 0
for i in cat:
    df[i] = df[i].replace('Yes',1)
    df[i] = df[i].replace('No',0)
```

```
In [24]: df
```

Out[24]:

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService
0	Female	0	1	0	1	0	0	
1	Male	0	0	0	34	1	0	
2	Male	0	0	0	2	1	0	
3	Male	0	0	0	45	0	0	
4	Female	0	0	0	2	1	0	Fiber
...	...	...	...	...	...	...	...	
7038	Male	0	1	1	24	1	1	
7039	Female	0	1	1	72	1	1	Fiber
7040	Female	0	1	1	11	0	0	
7041	Male	1	1	0	4	1	1	Fiber
7042	Male	0	0	0	66	1	0	Fiber

7043 rows × 20 columns

```
In [25]: from sklearn.preprocessing import OrdinalEncoder
oe = OrdinalEncoder()
#from sklearn.preprocessing import OrdinalEncoder

# Convert all values in 'InternetService' to strings
df['InternetService'] = df['InternetService'].astype(str)

# Create and fit the OrdinalEncoder
oe = OrdinalEncoder()
df['InternetService'] = oe.fit_transform(df[['InternetService']])

df['Contract'] = oe.fit_transform(df[['Contract']])
df['PaymentMethod'] = oe.fit_transform(df[['PaymentMethod']])
# df['Gender'] = oe.fit_transform(df['Gender'])
from sklearn.preprocessing import LabelEncoder
label_encoder = LabelEncoder()
df['gender'] = label_encoder.fit_transform(df['gender'])
df
```

```
Out[25]:
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService
0	0	0	1	0	1	0	0	
1	1	0	0	0	34	1	0	
2	1	0	0	0	2	1	0	
3	1	0	0	0	45	0	0	
4	0	0	0	0	2	1	0	
...	...	...	...	...	...	...	...	...
7038	1	0	1	1	24	1	1	
7039	0	0	1	1	72	1	1	
7040	0	0	1	1	11	0	0	
7041	1	1	1	0	4	1	1	
7042	1	0	0	0	66	1	0	

7043 rows × 20 columns

```
In [26]: scale_cols = ['tenure', 'MonthlyCharges', 'TotalCharges']
# now we scaling all the data
from sklearn.preprocessing import MinMaxScaler
scale = MinMaxScaler()
df[scale_cols] = scale.fit_transform(df[scale_cols])
df
```

Out[26]:

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	Internet
0	0	0	1	0	0.013889	0	0	
1	1	0	0	0	0.472222	1	0	
2	1	0	0	0	0.027778	1	0	
3	1	0	0	0	0.625000	0	0	
4	0	0	0	0	0.027778	1	0	
...	...	...	...	...	...	...	...	...
7038	1	0	1	1	0.333333	1	1	
7039	0	0	1	1	1.000000	1	1	
7040	0	0	1	1	0.152778	0	0	
7041	1	1	1	0	0.055556	1	1	
7042	1	0	0	0	0.916667	1	0	

7043 rows × 20 columns

In [27]: `df.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 20 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   gender                7043 non-null  int32  
 1   SeniorCitizen         7043 non-null  int64  
 2   Partner               7043 non-null  int64  
 3   Dependents            7043 non-null  int64  
 4   tenure                7043 non-null  float64
 5   PhoneService          7043 non-null  int64  
 6   MultipleLines         7043 non-null  int64  
 7   InternetService       7043 non-null  float64
 8   OnlineSecurity        7043 non-null  int64  
 9   OnlineBackup          7043 non-null  int64  
10  DeviceProtection      7043 non-null  int64  
11  TechSupport           7043 non-null  int64  
12  StreamingTV           7043 non-null  int64  
13  StreamingMovies       7043 non-null  int64  
14  Contract              7043 non-null  float64
15  PaperlessBilling      7043 non-null  int64  
16  PaymentMethod         7043 non-null  float64
17  MonthlyCharges        7043 non-null  float64
18  TotalCharges          7043 non-null  float64
19  Churn                 7043 non-null  int64  
dtypes: float64(6), int32(1), int64(13)
memory usage: 1.0 MB

```

## Independent and Dependent variables

```

In [28]: # independent and dependent variables
x = df.drop('Churn',axis=1)
y = df['Churn']

```

## Splitting data

```
In [29]: from sklearn.model_selection import train_test_split
xtrain,xtest,ytrain,ytest = train_test_split(x,y,test_size=0.2,random_state=10)
print(xtrain.shape)
print(xtest.shape)
```

(5634, 19)

(1409, 19)

## 1. Building Neural Network for Customer Churn Data

```
In [30]: # import tensorflow
import tensorflow as tf
#import keras
from tensorflow import keras
```

## Define Model

```
In [31]: # define se`quential model
model = keras.Sequential([
    # input layer
    keras.layers.Dense(19, input_shape=(19,), activation='relu'),
    keras.layers.Dense(15, activation='relu'),
    keras.layers.Dense(10,activation = 'relu'),
    # we use sigmoid for binary output
    # output layer
    keras.layers.Dense(1, activation='sigmoid')
])
```

```
In [32]: # xtrain = np.array(xtrain, dtype=np.float32) # Example: Convert xtrain to float32
# ytrain = np.array(ytrain, dtype=np.float32) # Example: Convert ytrain to float32
```

## Compile the Customer Churn Model

```
In [33]: # time for compilation of neural net.
model.compile(optimizer = 'adam',
              loss = 'binary_crossentropy',
              metrics = ['accuracy'])
# now we fit our model to training data
model.fit(xtrain,ytrain,epochs=100)
```

```
Epoch 1/100
177/177 [=====] - 3s 4ms/step - loss: 0.5090 - accuracy:
0.7428
Epoch 2/100
177/177 [=====] - 1s 3ms/step - loss: 0.4403 - accuracy:
0.7836
Epoch 3/100
177/177 [=====] - 1s 3ms/step - loss: 0.4340 - accuracy:
0.7943
Epoch 4/100
177/177 [=====] - 1s 3ms/step - loss: 0.4296 - accuracy:
0.7962
Epoch 5/100
177/177 [=====] - 1s 4ms/step - loss: 0.4287 - accuracy:
0.7962
Epoch 6/100
177/177 [=====] - 1s 3ms/step - loss: 0.4251 - accuracy:
0.8016
Epoch 7/100
177/177 [=====] - 1s 3ms/step - loss: 0.4240 - accuracy:
0.7985
Epoch 8/100
177/177 [=====] - 1s 3ms/step - loss: 0.4233 - accuracy:
0.8012
Epoch 9/100
177/177 [=====] - 1s 3ms/step - loss: 0.4199 - accuracy:
0.8055
Epoch 10/100
177/177 [=====] - 1s 3ms/step - loss: 0.4198 - accuracy:
0.8037
Epoch 11/100
177/177 [=====] - 1s 3ms/step - loss: 0.4195 - accuracy:
0.8032
Epoch 12/100
177/177 [=====] - 1s 3ms/step - loss: 0.4181 - accuracy:
0.8042
Epoch 13/100
177/177 [=====] - 1s 4ms/step - loss: 0.4175 - accuracy:
0.8035
Epoch 14/100
177/177 [=====] - 1s 4ms/step - loss: 0.4176 - accuracy:
0.8035
Epoch 15/100
177/177 [=====] - 1s 4ms/step - loss: 0.4155 - accuracy:
0.8076
Epoch 16/100
177/177 [=====] - 1s 4ms/step - loss: 0.4148 - accuracy:
0.8065
Epoch 17/100
177/177 [=====] - 1s 4ms/step - loss: 0.4136 - accuracy:
0.8048
Epoch 18/100
177/177 [=====] - 1s 4ms/step - loss: 0.4152 - accuracy:
0.8083
Epoch 19/100
177/177 [=====] - 1s 4ms/step - loss: 0.4156 - accuracy:
0.8076
Epoch 20/100
177/177 [=====] - 1s 4ms/step - loss: 0.4131 - accuracy:
0.8094
Epoch 21/100
177/177 [=====] - 1s 6ms/step - loss: 0.4128 - accuracy:
0.8071
Epoch 22/100
```

```
177/177 [=====] - 1s 6ms/step - loss: 0.4122 - accuracy:
0.8074
Epoch 23/100
177/177 [=====] - 1s 6ms/step - loss: 0.4116 - accuracy:
0.8074
Epoch 24/100
177/177 [=====] - 1s 7ms/step - loss: 0.4101 - accuracy:
0.8072
Epoch 25/100
177/177 [=====] - 1s 8ms/step - loss: 0.4094 - accuracy:
0.8092
Epoch 26/100
177/177 [=====] - 1s 7ms/step - loss: 0.4109 - accuracy:
0.8094
Epoch 27/100
177/177 [=====] - 1s 7ms/step - loss: 0.4105 - accuracy:
0.8062
Epoch 28/100
177/177 [=====] - 1s 7ms/step - loss: 0.4085 - accuracy:
0.8101
Epoch 29/100
177/177 [=====] - 1s 7ms/step - loss: 0.4079 - accuracy:
0.8099
Epoch 30/100
177/177 [=====] - 1s 7ms/step - loss: 0.4072 - accuracy:
0.8094
Epoch 31/100
177/177 [=====] - 1s 6ms/step - loss: 0.4078 - accuracy:
0.8115
Epoch 32/100
177/177 [=====] - 1s 6ms/step - loss: 0.4056 - accuracy:
0.8110
Epoch 33/100
177/177 [=====] - 1s 6ms/step - loss: 0.4060 - accuracy:
0.8117
Epoch 34/100
177/177 [=====] - 1s 6ms/step - loss: 0.4062 - accuracy:
0.8101
Epoch 35/100
177/177 [=====] - 1s 6ms/step - loss: 0.4074 - accuracy:
0.8081
Epoch 36/100
177/177 [=====] - 1s 6ms/step - loss: 0.4047 - accuracy:
0.8104
Epoch 37/100
177/177 [=====] - 1s 6ms/step - loss: 0.4046 - accuracy:
0.8110
Epoch 38/100
177/177 [=====] - 1s 6ms/step - loss: 0.4041 - accuracy:
0.8094
Epoch 39/100
177/177 [=====] - 1s 6ms/step - loss: 0.4041 - accuracy:
0.8140
Epoch 40/100
177/177 [=====] - 1s 6ms/step - loss: 0.4043 - accuracy:
0.8119
Epoch 41/100
177/177 [=====] - 1s 6ms/step - loss: 0.4033 - accuracy:
0.8108
Epoch 42/100
177/177 [=====] - 1s 6ms/step - loss: 0.4021 - accuracy:
0.8104
Epoch 43/100
177/177 [=====] - 1s 6ms/step - loss: 0.4022 - accuracy:
```

```
0.8127
Epoch 44/100
177/177 [=====] - 1s 6ms/step - loss: 0.4029 - accuracy:
0.8133
Epoch 45/100
177/177 [=====] - 1s 6ms/step - loss: 0.4004 - accuracy:
0.8149
Epoch 46/100
177/177 [=====] - 1s 4ms/step - loss: 0.4011 - accuracy:
0.8165
Epoch 47/100
177/177 [=====] - 1s 4ms/step - loss: 0.4018 - accuracy:
0.8135
Epoch 48/100
177/177 [=====] - 1s 3ms/step - loss: 0.3995 - accuracy:
0.8138
Epoch 49/100
177/177 [=====] - 1s 3ms/step - loss: 0.4002 - accuracy:
0.8136
Epoch 50/100
177/177 [=====] - 1s 3ms/step - loss: 0.3998 - accuracy:
0.8154
Epoch 51/100
177/177 [=====] - 1s 3ms/step - loss: 0.3994 - accuracy:
0.8159
Epoch 52/100
177/177 [=====] - 1s 3ms/step - loss: 0.3996 - accuracy:
0.8110
Epoch 53/100
177/177 [=====] - 1s 3ms/step - loss: 0.3980 - accuracy:
0.8165
Epoch 54/100
177/177 [=====] - 1s 4ms/step - loss: 0.3975 - accuracy:
0.8143
Epoch 55/100
177/177 [=====] - 1s 4ms/step - loss: 0.3963 - accuracy:
0.8133
Epoch 56/100
177/177 [=====] - 1s 7ms/step - loss: 0.3968 - accuracy:
0.8161
Epoch 57/100
177/177 [=====] - 1s 6ms/step - loss: 0.3962 - accuracy:
0.8166
Epoch 58/100
177/177 [=====] - 1s 7ms/step - loss: 0.3970 - accuracy:
0.8143
Epoch 59/100
177/177 [=====] - 1s 6ms/step - loss: 0.3957 - accuracy:
0.8131
Epoch 60/100
177/177 [=====] - 1s 6ms/step - loss: 0.3966 - accuracy:
0.8168
Epoch 61/100
177/177 [=====] - 1s 6ms/step - loss: 0.3959 - accuracy:
0.8136
Epoch 62/100
177/177 [=====] - 1s 7ms/step - loss: 0.3937 - accuracy:
0.8177
Epoch 63/100
177/177 [=====] - 1s 6ms/step - loss: 0.3942 - accuracy:
0.8174
Epoch 64/100
177/177 [=====] - 1s 5ms/step - loss: 0.3938 - accuracy:
0.8147
```



```
Epoch 65/100
177/177 [=====] - 1s 4ms/step - loss: 0.3945 - accuracy:
0.8158
Epoch 66/100
177/177 [=====] - 1s 4ms/step - loss: 0.3935 - accuracy:
0.8152
Epoch 67/100
177/177 [=====] - 1s 4ms/step - loss: 0.3939 - accuracy:
0.8172
Epoch 68/100
177/177 [=====] - 1s 4ms/step - loss: 0.3923 - accuracy:
0.8166
Epoch 69/100
177/177 [=====] - 1s 4ms/step - loss: 0.3925 - accuracy:
0.8163
Epoch 70/100
177/177 [=====] - 1s 4ms/step - loss: 0.3931 - accuracy:
0.8175
Epoch 71/100
177/177 [=====] - 1s 4ms/step - loss: 0.3926 - accuracy:
0.8197
Epoch 72/100
177/177 [=====] - 1s 4ms/step - loss: 0.3912 - accuracy:
0.8198
Epoch 73/100
177/177 [=====] - 1s 4ms/step - loss: 0.3922 - accuracy:
0.8154
Epoch 74/100
177/177 [=====] - 1s 4ms/step - loss: 0.3897 - accuracy:
0.8197
Epoch 75/100
177/177 [=====] - 1s 4ms/step - loss: 0.3921 - accuracy:
0.8143
Epoch 76/100
177/177 [=====] - 1s 4ms/step - loss: 0.3901 - accuracy:
0.8154
Epoch 77/100
177/177 [=====] - 1s 4ms/step - loss: 0.3897 - accuracy:
0.8182
Epoch 78/100
177/177 [=====] - 1s 4ms/step - loss: 0.3883 - accuracy:
0.8198
Epoch 79/100
177/177 [=====] - 1s 4ms/step - loss: 0.3883 - accuracy:
0.8195
Epoch 80/100
177/177 [=====] - 1s 4ms/step - loss: 0.3888 - accuracy:
0.8213
Epoch 81/100
177/177 [=====] - 1s 4ms/step - loss: 0.3895 - accuracy:
0.8163
Epoch 82/100
177/177 [=====] - 1s 4ms/step - loss: 0.3882 - accuracy:
0.8207
Epoch 83/100
177/177 [=====] - 1s 6ms/step - loss: 0.3870 - accuracy:
0.8248
Epoch 84/100
177/177 [=====] - 1s 6ms/step - loss: 0.3865 - accuracy:
0.8193
Epoch 85/100
177/177 [=====] - 1s 6ms/step - loss: 0.3866 - accuracy:
0.8172
Epoch 86/100
```

```

177/177 [=====] - 1s 8ms/step - loss: 0.3866 - accuracy:
0.8182
Epoch 87/100
177/177 [=====] - 1s 7ms/step - loss: 0.3869 - accuracy:
0.8204
Epoch 88/100
177/177 [=====] - 1s 7ms/step - loss: 0.3860 - accuracy:
0.8186
Epoch 89/100
177/177 [=====] - 1s 7ms/step - loss: 0.3865 - accuracy:
0.8166
Epoch 90/100
177/177 [=====] - 1s 7ms/step - loss: 0.3870 - accuracy:
0.8191
Epoch 91/100
177/177 [=====] - 1s 8ms/step - loss: 0.3844 - accuracy:
0.8202
Epoch 92/100
177/177 [=====] - 1s 7ms/step - loss: 0.3844 - accuracy:
0.8223
Epoch 93/100
177/177 [=====] - 1s 7ms/step - loss: 0.3849 - accuracy:
0.8227
Epoch 94/100
177/177 [=====] - 1s 7ms/step - loss: 0.3861 - accuracy:
0.8207
Epoch 95/100
177/177 [=====] - 1s 7ms/step - loss: 0.3840 - accuracy:
0.8232
Epoch 96/100
177/177 [=====] - 1s 7ms/step - loss: 0.3829 - accuracy:
0.8220
Epoch 97/100
177/177 [=====] - 1s 7ms/step - loss: 0.3839 - accuracy:
0.8220
Epoch 98/100
177/177 [=====] - 1s 7ms/step - loss: 0.3828 - accuracy:
0.8222
Epoch 99/100
177/177 [=====] - 1s 7ms/step - loss: 0.3821 - accuracy:
0.8248
Epoch 100/100
177/177 [=====] - 1s 7ms/step - loss: 0.3815 - accuracy:
0.8216

```

Out[33]: <keras.src.callbacks.History at 0x14ef7380790>

## evalute the model

In [34]: `model.evaluate(xtest,ytest)`

```

45/45 [=====] - 1s 3ms/step - loss: 0.4320 - accuracy: 0.
7800

```

Out[34]: [0.4320320785045624, 0.779985785484314]

## predict the churn values

In [35]: 

```
# predict the churn values
ypred = model.predict(xtest)
print(ypred)
# unscaling the ypred values
ypred_lis = []
```

```
for i in ypred:  
    if i>0.5:  
        ypred_lis.append(1)  
    else:  
        ypred_lis.append(0)  
print(ypred_lis)
```

45/45 [=====] - 0s 5ms/step

[[0.30121344]

[0.96312666]

[0.22262603]

...

[0.2961771 ]

[0.11197827]

[0.9920137 ]]

```

[0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0,
0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 1,
0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 1, 0, 0, 1, 0, 0, 0,
0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0,
1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0,
1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0,
0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0,
1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0,
0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1,
0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1,
0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 1, 1, 0, 0,
0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0,
0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0,
0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1,
1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0,
0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0,
1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0,
0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 1, 1,
0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 1, 0, 0, 0,
0, 1, 1, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1,
1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0,
0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0,
1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0,
0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1,
0, 0, 0, 0, 1]

```

```

In [36]: #make dataframe for comparing the original and predict values
data = {'original_churn':ytest, 'predicted_churn':ypred_lis}

```

```
df_check = pd.DataFrame(data)
df_check.head(10)
```

Out[36]:

	original_churn	predicted_churn
6418	0	0
1948	1	1
4497	0	0
66	0	0
1705	0	0
924	0	0
1051	0	0
7012	0	0
3723	0	0
4590	0	0

## Performance Matrices

In [37]:


```
# checking for performance metrics
# importing classification_report and confusion metrics
from sklearn.metrics import confusion_matrix, classification_report
# print classification_report
print(classification_report(ytest,ypred_lis))
```

	precision	recall	f1-score	support
0	0.85	0.86	0.86	1066
1	0.55	0.53	0.54	343
accuracy			0.78	1409
macro avg	0.70	0.70	0.70	1409
weighted avg	0.78	0.78	0.78	1409

In [38]:

```
# plotting the confusion matrix plot
conf_mat = confusion_matrix(ytest, ypred_lis)
# conf_mat = df.confusion_matrix(labels=ytest, predictions=ypred_lis)
plt.figure(figsize = (17,7))
sns.heatmap(conf_mat, annot=True,fmt='d')
plt.xlabel('Predicted_number')
plt.ylabel('True_number')
```

Out[38]: Text(183.2222222222223, 0.5, 'True\_number')



## Independent and Dependent

```
In [39]: x=df.drop(columns=['Churn'],axis=1)
         y=df['Churn']
```

```
In [40]: y.value_counts()# imbalance
```

```
Out[40]: Churn
0      5174
1      1869
Name: count, dtype: int64
```

```
In [41]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25,random_state=44)

print(f"the shape of x_train is : {x_train.shape}")
print(f'the shape of x_test is : {x_test.shape}')
print(f'the shape of y_train is : {y_train.shape}')
print(f'the shape of y_test is {y_test.shape}')
```

```
the shape of x_train is : (5282, 19)
the shape of x_test is : (1761, 19)
the shape of y_train is : (5282,)
the shape of y_test is (1761,)
```

## 2.Decision Tree Classifier


```
In [42]: from sklearn.tree import DecisionTreeClassifier
         dt=DecisionTreeClassifier(criterion="entropy", max_depth = 4)
         model2=dt.fit(x_train,y_train)
```

```
In [43]: y_pred=model2.predict(x_test)
         print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.82	0.90	0.86	1298
1	0.63	0.45	0.52	463
accuracy			0.78	1761
macro avg	0.72	0.68	0.69	1761
weighted avg	0.77	0.78	0.77	1761

```
In [44]: # plotting the confusion matrix plot
conf_mat = confusion_matrix(y_test, y_pred)
# conf_mat = df.confusion_matrix(labels=ytest, predictions=y_pred)
plt.figure(figsize = (17,7))
sns.heatmap(conf_mat, annot=True,fmt='d')
plt.xlabel('Predicted_number')
plt.ylabel('True_number')
```

Out[44]: Text(183.2222222222223, 0.5, 'True\_number')



### 3. Logistics Regression


```
In [45]: from sklearn.linear_model import LogisticRegression
lr=LogisticRegression()
model3=lr.fit(x_train,y_train)
```

```
In [46]: y_pred=model3.predict(x_test)
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.84	0.90	0.87	1298
1	0.65	0.53	0.58	463
accuracy			0.80	1761
macro avg	0.74	0.71	0.72	1761
weighted avg	0.79	0.80	0.79	1761

```
In [47]: # plotting the confusion matrix plot
conf_mat = confusion_matrix(y_test, y_pred)
# conf_mat = df.confusion_matrix(labels=ytest, predictions=y_pred)
plt.figure(figsize = (17,7))
sns.heatmap(conf_mat, annot=True,fmt='d')
plt.xlabel('Predicted_number')
plt.ylabel('True_number')
```

Out[47]: Text(183.2222222222223, 0.5, 'True\_number')



## 4.Random Forest Classifier

```
In [48]: from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25,random_state=44)
```

```
In [49]: # Create a Random Forest Classifier
rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)

# Fit the model to the resampled training data
rf_classifier.fit(x_train, y_train)

# Make predictions on the test data
y_pred = rf_classifier.predict(x_test)
```


```
In [50]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.82	0.89	0.86	1298
1	0.60	0.45	0.51	463
accuracy			0.78	1761
macro avg	0.71	0.67	0.68	1761
weighted avg	0.76	0.78	0.77	1761

```
In [51]: # plotting the confusion matrix plot
conf_mat = confusion_matrix(y_test, y_pred)
# conf_mat = df.confusion_matrix(labels=ytest,predictions=y_pred_lis)
plt.figure(figsize = (17,7))
sns.heatmap(conf_mat, annot=True,fmt='d')
plt.xlabel('Predicted_number')
plt.ylabel('True_number')
```

```
Out[51]: Text(183.2222222222223, 0.5, 'True_number')
```





## 5. SVM

```
In [52]: from sklearn.svm import SVC
from sklearn.metrics import classification_report
```

```
In [53]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25,random_state=44)
# Create an SVM classifier
svm_classifier = SVC(kernel='linear', C=1, random_state=42)

# Fit the SVM model to the training data
svm_classifier.fit(x_train, y_train)
```

```
Out[53]: SVC
SVC(C=1, kernel='linear', random_state=42)
```

```
In [54]: # Make predictions on the test data
y_pred_svm = svm_classifier.predict(x_test)

# Generate a classification report for the SVM model
svm_classification_report = classification_report(y_test, y_pred_svm)
```

```
In [55]: # Print the classification report for the SVM model
print("SVM Classifier Classification Report:")
print(svm_classification_report)
```

```
SVM Classifier Classification Report:
              precision    recall  f1-score   support

     0       0.85         0.89         0.87         1298
     1       0.65         0.55         0.59          463

 accuracy          0.80         0.80         0.80         1761
 macro avg         0.75         0.72         0.73         1761
 weighted avg         0.80         0.80         0.80         1761
```

```
In [56]: # plotting the confusion matrix plot
conf_mat = confusion_matrix(y_test, y_pred)
# conf_mat = df.confusion_matrix(labels=ytest,predictions=y_pred_lis)
plt.figure(figsize = (17,7))
sns.heatmap(conf_mat, annot=True,fmt='d')
```

```
plt.xlabel('Predicted_number')  
plt.ylabel('True_number')
```

Out[56]: Text(183.22222222222223, 0.5, 'True\_number')

